

NewzTrader: Autonomous Trading Agent Implementation Using Natural Language Processing Of News Headlines

William Lyon

November 18, 2012

Abstract

Natural Language Processing techniques are used to examine financial news headlines and generate predictions of stock price movements. News headlines from the Wall Street Journal from January 1, 2009 to November 1, 2012 are collected and paired with daily S&P 500 index returns. This information is used to train a Naive Bayes classifier and generate BUY/SELL trading signals. This trading strategy is then backtested using Wall Street Journal news headlines as a predictor for movement in the price of the S&P 500 index.




Figure 1: Cumulative frequency plot

1 Introduction

In finance, the Efficient Market Hypothesis states that all publicly available information is reflected in financial market prices. As new information becomes available, prices adjust to take this new information into account.

1.1 Motivation

This tool could be used as a component of an autonomous trading agent that will make BUY/SELL decisions for trading financial instruments.

Goals: 1)

1.2 Literature Review

2 Examining The Data

Google Finance provides access to historical stock quotes as well as links to financial news stories pertaining to specific stocks. This data provides the basis for this paper.

2.1 Historical Stock Market Quotes

2.2 News Headlines

A total of 101,618 news headlines over 1416 days were collected with an average of 71 headlines per day. After aligning with financial data: UP: 14651 DOWN: 13718 NONE: 57405 TO-TOTAL: 85775

Prior probs: UP: 0.1708 DOWN: 0.1599 NONE: 0.6693

3 Classification Methodology

These tuples of [return, listOfDailyNews] are used as the training data for a Naive Bayes classifier.

3.1 Classes

3.2 Naive Bayes

3.3 Maximum Entropy

4 Implementation

4.1 Dependencies

4.2 Data Collection & Munging

4.3 Feature Extraction

4.3.1 Bag of words

4.3.2 Filtering stopwords

4.3.3 Include significant bigrams

4.4 Naive Bayes Classifier

4.5 Maximum Entropy Classifier

4.6 Backtesting

5 Evaluation

5.1 Accuracy

5.2 Most informative features

5.3 Precision

5.4 Recall

6 Trading Model

6.1 Trading Signals

6.2 Backtesting

9 Source Code Listings

Listing 1: newsCredScraper.py: scrapes WSJ news headlines

```
1
2 from lxml import etree
3 from datetime import datetime
4 from datetime import date
5 from datetime import timedelta
6 #import datetime
7 from dateutil.parser import parse
8 #import pandas as pd
9 import pickle
10 #from pandas.io.data import DataReader
11 #from pandas.io.data import DataReader
12 #from pandas.io.data import DataReader
13
14 #import pandas as pd
15 #from datetime import datetime
16
17 #import pandas as pd
18 #
19 #from pandas.io.data import DataReader
20
21 def daterange(start_date, end_date):
22     for n in range(int((end_date - start_date).days)):
23         yield start_date + timedelta(n)
24
25 def only_alphanumeric(s):
26     #s = unicode(s, "utf-8")
27     return ''.join(c for c in s.split() if c.isalnum())
28 def only_alpha(s):
29     return ''.join(c for c in s.split() if c.isalpha())
30 def removeNonAscii(s): return ''.join(i for i in s if ord(i)<128)
31
32 #sp500 = DataReader("SPY", "yahoo", datetime(2009, 1, 1))
33 news={}
34 #descs = {}
35 start_date=date(2011, 7, 27)
36 end_date = date(2012, 11, 9)
37 news[date(2009, 11, 9)] = []
38 news[date(2009, 11, 10)] = []
39 for single_date in daterange(start_date, end_date):
40     news[single_date]=[]
41 for single_date in daterange(start_date, end_date):
42     #print "Starting date: " + str(single_date)
43     year = single_date.year
44     month = single_date.month
45     day = single_date.day
46     path = "http://api.newscred.com/articles?access_key=c4bcc3f7c9bf9ec159f51da0a86ca658&sources=104
         afa30d811d37a5582a39e1662a311&pagesize=99&from_date=%d-%d-%d&to_date=%d-%d-%d_23:59:59" % (year,
         month, day, year, month, day)
47     while True:
48         try:
49             root = etree.parse(path)
50             break
51         except etree.XMLSyntaxError:
52             pass
53     myRoot = root.getroot()
54
55     #news[date(year, month, day)]=[]
56     #descs[date(year, month, day)]=[]
```

```

57     for element in myRoot.iter("article"):
58         #for item in element.iter("description"):
59             # desc = item.text
60         for item in element.iter("title"):
61             title = item.text
62         for item in element.iter("created_at"):
63             pubDate = parse(item.text)
64
65         news[pubDate.date()].append(only_alphanum(removeNonAscii(title)))
66 #import pickle
67     output = open('newsDict2.pkl', 'wb')
68     pickle.dump(news, output)
69     output.close()

```

Listing 2: dataGetter.py: downloads historical S&P 500 index price data and joins with WSJ news headlines / identifies news headline classifications and saves in format for corpus reader

```

1  # dataGetter.py
2  # William Lyon
3  # AI Grad Project
4  # NewzTrader
5  # dataGetter.py
6  # 1) Loads pickled news dict
7  # 2) Downloads historical stock price data
8  # 3) Joins news dict and stock data in a pandas dataframe
9  # 4) Set UP/DOWN/NONE classifications for every trading day
10 # 5) Save news headlines in .CSV files for corpus reader
11
12 import pickle
13 from datetime import datetime
14 from datetime import timedelta
15 from datetime import date
16
17 def daterange(start_date, end_date):
18     for n in range(int ((end_date - start_date).days)):
19         yield start_date + timedelta(n)
20
21 fkl_file = open('combinedNewsDictFull.pkl', 'rb')
22 news = pickle.load(fkl_file)
23 fkl_file.close()
24 start_date=date(2009, 1, 1)
25 end_date = date(2012, 11, 17)
26
27 import pandas as pd
28 #
29 from pandas.io.data import DataReader
30
31 def only_alphanum(s):
32     #s = unicode(s, "utf-8")
33     return ''.join(c for c in s.split() if c.isalnum())
34 def only_alpha(s):
35     return ''.join(c for c in s.split() if c.isalpha())
36 def removeNonAscii(s): return ''.join(i for i in s if ord(i)<128)
37
38 sp500 = DataReader("^GSPC", "yahoo", datetime(2009, 1, 1))
39 newsframe = pd.Series(news, name='News')
40 #descframe = pd.Series(descs, name='Desc')
41 frameWithNews = sp500.join(pd.DataFrame(newsframe))
42 #sp500Frame = frameWithNews.join(pd.DataFrame(descframe))
43 newsframe = pd.Series(news, name='News')
44 #descframe = pd.Series(descs, name='Desc')
45 frameWithNews = sp500.join(pd.DataFrame(newsframe))

```

```

46 newsReturns = frameWithNews['Adj_Close'].pct_change()
47 newsReturns.name='Returns'
48 returnsFrame = frameWithNews.join(pd.DataFrame(newsReturns))
49
50
51 returnsFrame['UP'] = returnsFrame>Returns > 0.01
52 returnsFrame['DOWN'] = returnsFrame>Returns < -0.01
53 returnsFrame['NONE'] = (returnsFrame['UP']==False) & (returnsFrame['DOWN']==False)
54
55 droppedFrame = returnsFrame
56
57 newsUP_frame = droppedFrame[droppedFrame['UP']==True]
58
59
60 newsDOWN_frame = droppedFrame[droppedFrame['DOWN']==True]
61 newsNONE_frame = droppedFrame[droppedFrame['NONE']==True]
62
63 newsNONE_frame = newsNONE_frame.dropna()
64 newsUP_frame = newsUP_frame.dropna()
65 newsDOWN_frame = newsDOWN_frame.dropna()
66
67 # DO THIS FOR NONE, UP, and DOWN
68 i = 1
69 for row in newsNONE_frame.iterrows():
70     i+=1
71     if len(row[1].ix['News'])>0:
72         for line in row[1].ix['News']:
73             i+=1
74             writeFile = open('%d_news_NONE.csv' % i, 'w')
75             writeFile.write(line+'\n')
76
77 # DO THIS FOR NONE, UP, and DOWN
78 i = 1
79 for row in newsUP_frame.iterrows():
80     i+=1
81     if len(row[1].ix['News'])>0:
82         for line in row[1].ix['News']:
83             i+=1
84             writeFile = open('%d_news_UP.csv' % i, 'w')
85             writeFile.write(line+'\n')
86
87 # DO THIS FOR NONE, UP, and DOWN
88 i = 1
89 for row in newsDOWN_frame.iterrows():
90     i+=1
91     if len(row[1].ix['News'])>0:
92         for line in row[1].ix['News']:
93             i+=1
94             writeFile = open('%d_news_DOWN.csv' % i, 'w')
95             writeFile.write(line+'\n')

```

Listing 3: nbTrainer.py: loads news corpus / trains Naive Bayes classifier

```

1 # nbTrainer.py
2 # William Lyon
3 # AI Grad Project
4 # NewzTrader
5 # nbTrainer.py
6 # 1) Load NLTK corpus reader for naive bayes classifier
7 # 2) Train NB classifier with random 90% of corpus features
8 # 3) Test NB classifier with remaining 10% of corpus features and report
9 # accuracy

```

```

10 # TODO: recall, precision reports; improve classifier (only strong words?)
11
12 # NLTK – train nb_classifier
13
14
15 import random
16 import nltk as nltk
17 #nltk.download()
18 from nltk.corpus import stopwords
19 import os, os.path
20 path = os.path.expanduser('~/.nltk_data')
21 if not os.path.exists(path):
22     os.mkdir(path)
23 os.path.exists(path)
24 import nltk.data
25 path in nltk.data.path
26 from nltk.corpus.reader import CategorizedPlaintextCorpusReader
27 reader = CategorizedPlaintextCorpusReader('!', r'.*_news_.*\.csv', cat_pattern=r'.*_news_(\w+)\.csv')
28 reader.categories()
29
30 def bag_of_words(words):
31     return dict([(word, True) for word in words if word[0].isalpha()])
32 import collections
33 def bag_of_words_not_in_set(words, badwords):
34     return bag_of_words(set(words) - set(badwords))
35
36 def bag_of_non_stopwords(words, stopfile='english'):
37     badwords = stopwords.words(stopfile)
38     return bag_of_words_not_in_set(words, badwords)
39
40 from nltk.metrics import BigramAssocMeasures
41 from nltk.collocations import BigramCollocationFinder
42
43 def bag_of_bigrams_words(words, score_fn=BigramAssocMeasures.chi_sq, n=2000):
44     bigram_finder = BigramCollocationFinder.from_words(words)
45     bigrams = bigram_finder.nbest(score_fn, n)
46     dictOfBigrams = bag_of_words(bigrams)
47     dictOfBigrams.update(bag_of_non_stopwords(words))
48     return dictOfBigrams
49
50 def label_feats_from_corpus(corp, feature_detector=bag_of_bigrams_words):
51     label_feats = collections.defaultdict(list)
52     for label in corp.categories():
53         for fileid in corp.fileids(categories=[label]):
54             feats = feature_detector(corp.words(fileids=[fileid]))
55             label_feats[label].append(feats)
56     return label_feats
57
58 def split_label_feats(lfeats, split=0.90):
59     train_feats = []
60     test_feats = []
61     for label, feats in lfeats.iteritems():
62         random.shuffle(feats, random.random)
63         cutoff = int(len(feats) * split)
64         train_feats.extend([(feat, label) for feat in feats[:cutoff]])
65         test_feats.extend([(feat, label) for feat in feats[cutoff:]])
66     return train_feats, test_feats
67
68
69
70

```

```
71 reader.categories()
72
73 lfeats = label_feats_from_corpus(reader)
74 lfeats.keys()
75 train_feats, test_feats = split_label_feats(lfeats)
76 len(train_feats)
77 len(test_feats)
78
79 from nltk.classify import NaiveBayesClassifier
80 nb_classifier = NaiveBayesClassifier.train(train_feats)
81 nb_classifier.labels()
82
83 from nltk.classify.util import accuracy
84 accuracy(nb_classifier, test_feats)
```

Listing 4: backTest.py: simulates trading using trading signals generated from WSJ news headlines using Naive Bayes classifier

Listing 5: NewzTrader.py